Translational AI Center

Module 4 — Communication & Leadership

Aditya Balu

5/8/2026

Recap — what you now know

Three modules of strategist preparation. One module to make it real.

- **Module 1 — Foundations —** mindset shift, lifecycle, roles, opportunity matrix, readiness
- **Module 2 — Landscape —** API tools, agentic AI, multi-modal, VLAs, 3D, GNNs
- **Module 3 — Risk —** NIST AI RMF, technical risks, governance, mitigation playbook
- **Module 4 — today —** the leadership and communication playbook to take all of this back home

From strategist to champion

Knowing what to do is half the job. Getting the org to do it is the other half.

Strategist

- **Knows** — the AI map and the risks
- **Picks** — the right tool for the right job
- **Documents** — the strategy
- **Limits** — without buy-in, nothing ships

Champion

- **Translates** — to every audience that matters
- **Recruits** — sponsors, allies, first followers
- **Manages** — fear, change, vendor relationships
- **Result** — AI strategy that actually leaves your laptop

Translational AI Center

Know Your Audience

Most AI initiatives fail at the people layer, not the tech.

Why people, not pipelines, kill most AI initiatives

The model worked. The change didn't.

- **The pilot succeeded — but no one used it —** no one was asked what they actually needed
- **Engineering shipped — but legal blocked launch —** compliance was looped in too late
- **Demo blew minds — but the sponsor moved on —** no one re-recruited the executive
- **Tool rolled out — but adoption sat at 8% —** no one trained or re-designed the workflow
- **Champion burned out —** lone effort with no coalition

Five stakeholder archetypes

Each one cares about different things. Speak to each in their language.

Executives

ROI, strategic positioning, peer benchmarks

Board / oversight

risk, reputation, fiduciary duty, compliance

Peers

fairness, turf, credit, parallel pilots

Teams

“does this make my job easier or harder?”

Customers

value, trust, transparency, control

Influence × Interest map

Where each stakeholder sits decides how you spend your time.

↑ Interest

Manage closely

high influence · high interest
Key players

Keep informed

low influence · high interest
Allies & advocates

Keep satisfied

high influence · low interest
Don't let them turn against you

Monitor

low influence · low interest
Light touch

Influence →

What each archetype actually wants to hear

The same initiative, five different headlines.

Executives	“Revenue impact, peer benchmarks, two-year position.”
Board / oversight	“Risk envelope, compliance posture, governance evidence.”
Peers	“What this means for your area, where credit is shared.”
Teams	“What changes for you on Monday morning. What stays the same.”
Customers	“What you get, what we promise, what we won't do with your data.”

EXERCISE — Map your stakeholders

5 minutes individual. Use the workbook page.

- **Pick one initiative —** from your opportunity matrix in Module 1
- **List 5–8 stakeholders —** by name, not job title
- **Place each —** on the influence × interest grid
- **Mark the gaps —** where you don't actually know what someone wants
- **Pick your first conversation —** the highest-influence person you've under-engaged

Translational AI Center

Speak Their Language

The same initiative, retold five different ways.

Translate value to each audience

Pick the dimension that matters to the person in front of you.

Finance

ROI · payback period · margin impact

Legal

risk envelope · regulatory exposure · contracts

Operations

throughput · error rate · capacity unlocked

Board

growth optionality · competitive position

Customers

experience · trust · transparency

Employees

augmentation · time saved · skills gained

The before / after frame

The fastest way to make value concrete.

BEFORE

- **Workflow** — 5 hours per claim, manual review
- **Errors** — 2-4% slip through to payments
- **Capacity** — 12 claims per day per FTE
- **Customer** — average 7 days to resolution

AFTER

- **Workflow** — AI extracts + flags edge cases for review
- **Errors** — 0.7%, with explainable flags
- **Capacity** — 30 claims per day per FTE
- **Customer** — average 1.5 days to resolution

Three story patterns that win

Pick one. Make it concrete. Tell it the same way every time.

1

The Hero

“We had a problem. Here's how we solved it. Here's what changed.”

2

The Contrast

“Old way vs new way. Same task, very different outcome.”

3

The Prediction

“This is the change coming. Here's how we'll be ready before competitors.”

Three concrete success narratives

Real templates you can adapt for your own initiatives.

- **Document extraction (operations) —** Insurance ops team cut average claim handling from 7 days to 1.5 — same headcount, more throughput, better customer NPS.
- **Knowledge assistant (employee) —** Internal RAG over manuals + SOPs cut new-hire ramp time from 90 days to 45. Tier-1 support tickets dropped 30%.
- **Predictive maintenance (asset-heavy) —** Vibration + heat-sensor model flagged a bearing failure 11 days early. One avoided shutdown paid for the year's AI program.

Notice — every one is concrete numbers, named workflow, and the outcome a non-technical leader cares about.

Anti-patterns — what NOT to say

If your message contains these, rewrite before sending.

Hype

“Revolutionary AI-powered next-generation solution…”

Jargon

“Leveraging transformer-based latent embeddings…”

Doom

“If we don't act now we'll be obsolete in 18 months.”

Vague

“Significant productivity gains across the enterprise.”

Antidote: specific numbers, named workflows, plain words.

Quick rewrite — fix this announcement

3 minutes. Discuss with a tablemate.

BEFORE (the hype version)

"We are thrilled to announce our revolutionary AI-powered next-generation platform that will transform how our team approaches every workflow, leveraging cutting-edge generative intelligence to unlock unprecedented productivity gains across the entire enterprise."

AFTER (your rewrite)

- What is this for? (one specific workflow)
- Who's the user? (one specific team)
- What changes for them on Monday?
- What's the measurable outcome we expect?
- What's NOT changing?

Translational AI Center

Fear Is Information

Resistance tells you what to fix, not who to fight.

The five most common AI fears

Chances are at least three of these are sitting in your audience right now.

Job loss

"Will AI replace me?"

Deskilling

"Will I lose what makes my work valuable?"

Bias / unfairness

"Will it treat me / my customers unfairly?"

Security

"Will my data, IP, secrets leak?"

Loss of human touch

"Will customers lose what they liked about us?"

The empathy template

Three lines you can use, in order, in any tense conversation about AI.

1

Acknowledge

“This change is real, and it's worth being concerned about.”

2

Specify

“Here's exactly what's changing — and exactly what isn't.”

3

Commit

“Here's the support you'll get. Here's how you can flag if it isn't working.”

Reframe — augmentation, not replacement

People who use AI well will replace people who don't. Help your team be the former.

- **The replacement framing —** “AI will do your job.” — fear, defensiveness, resistance
- **The augmentation framing —** “AI will do the boring 30% of your job. Your judgment is the other 70%.”
- **Why it works —** honors what people already do well — and asks them to do MORE of it
- **Where it fails —** if the augmentation framing isn't true. Don't say it if you mean replacement.
- **Rule —** tell the truth about what's changing. The empathy template handles the hard truths.

The literacy → fluency → strategist arc

The single most effective anti-fear narrative you have.

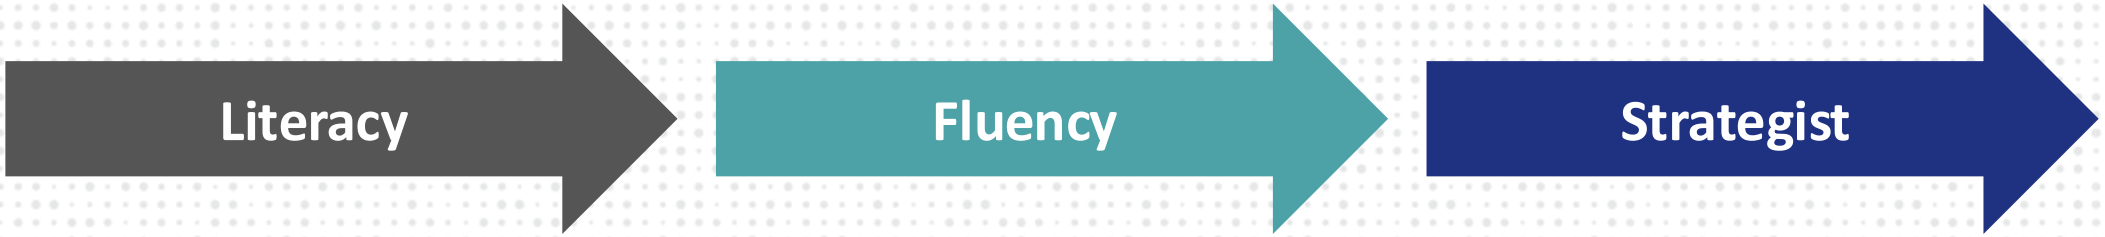

I know what AI is

I can prompt and get useful output

I know which family to use, when, and why

Tell every employee they're on this path. Make the path concrete. Provide the time and tools to walk it.

Translational AI Center

Leading the Change

Tools age. Models change. The change-management muscle is the one that lasts.

Why AI change management is different

Most change-management playbooks assume the system, once shipped, stays still.

- **AI keeps changing —** underlying models update; behavior shifts mid-flight
- **The user is also the test —** people learn what works through trial and error
- **Outcomes are probabilistic —** no two runs identical — that's hard for procedure-trained orgs
- **Risk profile is dynamic —** what was safe at 100 users may not be at 10,000
- **Implication —** communication can't be one-shot. It's a continuing rhythm.

The four stages of AI adoption

Each one needs a different communication strategy.

PILOT	ADOPT	SCALE	EMBED
narrow scope, real users, real metrics	early users champion it; expand to neighbors	platform thinking, training, governance	AI is part of the operating model, not a project

Most orgs get stuck between Adopt and Scale — the “demo graveyard.”

The trough of disillusionment

Excitement → frustration → quiet productivity. Most orgs quit in the middle.

- **Phase 1 — Excitement —** “This will change everything.” Pilots get green-lit. Early demos amaze.
- **Phase 2 — Frustration —** “This is harder than the demo.” Edge cases pile up. Costs are higher.
- **Phase 3 — Quiet productivity —** “Oh — this is actually working.” Less buzz, more value.
- **Champion's job —** narrate the trough as expected, not as failure
- **Failure mode —** executives quit at the trough — and miss phase 3 entirely

The champion's five moves

What you do every quarter to keep AI moving inside the org.

1 Tell stories

concrete wins, named teams, small numbers

2 Surface losses

what didn't work, what we learned, what we changed

3 Refresh sponsorship

executive's calendar moves; re-recruit quarterly

4 Recruit new champions

find the next-generation in adjacent teams

5 Update the strategy

the landscape changes; your strategy should too

Translational AI Center

Tech-to-Business Translation

The cheat sheet, the right questions, and when to push back.

Essential AI vocabulary — strategist's cheat sheet

The fifteen terms you'll hear most often. One sentence each.

Term	What it actually means
LLM	a model that predicts the next token in text
Embedding	a numeric vector that represents meaning
Token	a chunk of text the model processes (roughly word-piece)
Context window	the most text the model can read in one shot
RAG	retrieve relevant chunks, then feed them to the model
Fine-tune	further training on your data to specialize a model
Prompt	the instructions you give the model
Eval	an automated test that scores model output
Drift	model accuracy degrades as the world changes
Hallucination	plausible-but-false output
Inference	running the model to produce an output
Latency	time from request to response
Guardrail	a check on what the model can say or do
Agent	an LLM connected to tools that can take actions
MLOps	the discipline of running models in production

Translate both directions

Tech-to-business and business-to-tech are equally important.

Tech → Business

- **Engineer says** — “The eval F1 dropped 3 points after the model upgrade.”
- **You translate** — “We hit a quality regression. Customers may notice. Roll back or accept?”

Business → Tech

- **Exec says** — “We need this to be fair.”
- **You translate** — “Pick a fairness metric. Test the model against protected groups. Document the trade-off.”

Seven questions for any AI technical review

If your team can't answer these, it's not ready to ship.

- 1. What problem is this solving, in business language?
- 2. What does success look like — measured how, by whom?
- 3. What's the failure mode that ends up on the front page?
- 4. What's our human-in-the-loop pattern?
- 5. What's our cost per outcome — and our cost ceiling?
- 6. How do we know it's still working in 6 months?
- 7. What does “shut it off in 60 seconds” look like?

When and how to push back on engineers

Curiosity beats authority — and works better.

- **When the answer is too confident —** “What would have to be true for this NOT to work?”
- **When the answer is too vague —** “Show me the eval. Show me the dashboard.”
- **When the timeline shrinks suspiciously —** “What did we just remove from scope?”
- **When risk is dismissed —** “What's the worst case for an external observer?”
- **When you don't understand —** say so. Ask. The one who admits it first wins the room.

Translational AI Center

Working with Vendors

Most of your AI stack is going to be someone else's. Pick well, hold them accountable.

Build · Buy · Partner — a leadership lens

The default is buy. The exception is build. Partner is for the in-between.

Build	Buy	Partner
internal differentiation · sensitive data · regulated	commodity capability · fast time to market · supported	co-developed capability · shared risk · longer horizon

Tip — most decisions look like “build” to engineers and “buy” to finance. Map the trade-offs explicitly.

The mock vendor meeting

Three things to ask. Three things to avoid.

ASK

- **References** — two customers in our industry, on a call, this week
- **Eval** — let me run my evaluation on a new release before it ships
- **Data** — where it's processed, who can see it, retention policy
- **Roadmap** — what features are us paying you to defer?
- **Termination** — the off-ramp — data, weights, contracts

AVOID

- **Demo dazzle** — the canned demo is not your workflow
- **Logo lists** — “Fortune 500 customers” without your-industry references
- **Roadmap promises** — in writing or it didn't happen
- **Vague pricing** — “contact us” at $1M ARR is a red flag
- **No exit plan** — you'll regret signing without one

Negotiating an AI vendor contract — the red lines

Walk away if any of these aren't addressed.

- **Training on your data —** explicit no-train clause, default OFF, opt-out documented
- **Sub-processors —** list, change notice, audit rights
- **Model deprecation —** minimum notice, downgrade path, parallel-run window
- **Eval access —** you run your harness on candidate releases pre-GA
- **Incident response —** SLA, who pays for what, breach notification
- **Termination —** data return, deletion, transition assistance

After the contract — what to monitor

Vendor risk doesn't stop at signature.

- **Performance —** your eval harness on every release; trend, not single-point
- **Cost —** actual unit economics; alert on drift from contracted rates
- **Capability changes —** behavioral diffs across versions even when tests still pass
- **Regulatory posture —** SOC 2, ISO 27001, ISO 42001 — request annual updates
- **Strategic health —** vendor finances, ownership, leadership stability
- **Switchability —** could we leave in 90 days if we had to?

Translational AI Center

You Don't Lead Alone

AI initiatives die without coalitions.

Who to recruit — and why each one matters

Build the smallest coalition that can carry the change.

Sponsor

executive cover, budget, removes blockers

Ally

another champion in a different team

Skeptic

principled critic — turns into your fiercest advocate when convinced

First follower

early adopter who validates the change publicly

Critic-of-critics

someone who pushes back when AI hype goes too far — keeps you honest

The first follower

Your second supporter matters more than your first. Here's why.

- **The lone champion is fragile —** easy to dismiss as one enthusiast
- **The first follower legitimizes you —** now there are two — others can join
- **Find them early —** the most credible peer who's curious — not the most senior
- **Make them visible —** give them a podium moment, not a slide credit
- **Treat them like a partner —** their wins are your wins; their concerns are your roadmap

Communicate wins AND losses — early

Coalitions are built on honesty, not curated PR.

Wins

- **Concrete** — named team, named workflow, real numbers
- **Shared** — credit the team and the first follower, not yourself
- **Frequent** — small wins matter more than big launches
- **Repeatable** — “here's how someone else can do this too”

Losses

- **Honest** — what didn't work, what surprised us
- **Specific** — what we changed, what we'll watch next
- **Early** — before the rumor mill gets there
- **Re-positioning** — loss as data, not as defeat

Translational AI Center

The First 90 Days

What to do Monday morning and the eleven Mondays after.

Days 1–30 — Listen · Map · Prioritize

Resist the urge to ship anything yet.

- **Week 1 —** 1:1s with 10 stakeholders. Open-ended. What's frustrating? What works?
- **Week 2 —** use-case inventory. Map every AI experiment + tool already in use (shadow AI counts).
- **Week 3 —** data + readiness audit. Score against the four dimensions from Module 1.
- **Week 4 —** draft the opportunity matrix. Pick 2 quick wins + 1 strategic bet.
- **End of month —** a one-page strategy document. Share, don't publish.

Days 31–60 — Pilot · Learn · Measure

Now you ship. Small. With instruments.

- **Week 5 —** kick off the first pilot. Named owner. Named users. Named metrics.
- **Week 6 —** stand up minimum governance — risk register, kill switch, dashboard
- **Week 7 —** first real metric reading. Compare to the pre-pilot baseline.
- **Week 8 —** share the early result — wins AND surprises — to your coalition
- **End of month —** a go / iterate / kill decision on the pilot, in writing.

Days 61–90 — Scale · Communicate · Set the rhythm

Move from “a thing happening” to “the way we work.”

- **Week 9 —** expand the pilot to a second team. Apply the lessons.
- **Week 10 —** onboard new champions; share a written runbook
- **Week 11 —** first formal review with the AI Governance Committee
- **Week 12 —** publish externally — to customers, partners, the board — what you've learned
- **End of quarter —** set the cadence — monthly review, quarterly strategy refresh

Leading indicators of success

Before the headline ROI lands, watch these.

- **Time-to-pilot —** from idea to first user in days, not quarters
- **Adoption —** % of intended users actively using it
- **Quality —** eval scores, user feedback, escalation rate
- **Risk register growth + closure —** you're finding risks AND closing them
- **Coalition expansion —** new champions, new pilots in adjacent teams
- **Sponsor engagement —** they ask about it without you initiating

Translational AI Center

Pulling It All Together

The last quarter-mile of the workshop.

The value dashboard you'll build

Six metrics that prove AI is working — and warn you when it isn't.

Adoption

active users / intended users

Quality

eval scores, escalation rate

Outcome

the business KPI tied to the use case

Cost

$/outcome — and trend

Risk

open risks, closed risks, severity

Coalition

active champions, pilots in flight

Continuous learning paths

What to read, listen to, and follow — generic categories that age well.

- **Industry research —** any vendor-neutral research firm covering enterprise AI
- **Practitioner blogs —** engineering blogs from AI labs and AI-native companies
- **Policy + governance —** your local government's AI policy office; NIST AI RMF updates
- **Communities —** local AI meetups, professional associations in your sector
- **Hands-on —** spend two hours a week prompting, building, evaluating — don't outsource literacy
- **Internal —** your engineering team's eval reports — the most valuable read in your inbox

Three takeaways from the day

1

LLMs are one tool

You now have a mental map of the whole AI landscape — agentic, multi-modal, VLAs, 3D, GNNs.

2

Risk is your job

NIST AI RMF anchors the work. Govern · Map · Measure · Manage, every quarter.

3

Lead with stories

Strategy you can't communicate doesn't ship. Concrete numbers, named workflows, plain words.